

linear_regression

January 23, 2024

0.1 Linear regression workbook

This workbook will walk you through a linear regression example. It will provide familiarity with Jupyter Notebook and Python. Please print (to pdf) a completed version of this workbook for submission with HW #1.

ECE C147/C247, Winter Quarter 2024, Prof. J.C. Kao, TAs: T.Monsoor, Y. Liu, S. Rajesh, L. Julakanti, K. Pang

```
[ ]: import numpy as np
import matplotlib.pyplot as plt

#allows matlab plots to be generated in line
%matplotlib inline
```


0.1.1 Data generation

For any example, we first have to generate some appropriate data to use. The following cell generates data according to the model: $y = x - 2x^2 + x^3 + \epsilon$

```
[ ]: np.random.seed(0) # Sets the random seed.
num_train = 200 # Number of training data points

# Generate the training data
x = np.random.uniform(low=0, high=1, size=(num_train,))
y = x - 2*x**2 + x**3 + np.random.normal(loc=0, scale=0.03, size=(num_train,))
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')
```

```
[ ]: Text(0, 0.5, '$y$')
```



0.1.2 QUESTIONS:

Write your answers in the markdown cell below this one:

- (1) What is the generating distribution of x ?
- (2) What is the distribution of the additive noise ϵ ?

0.1.3 ANSWERS:

- (1) Based upon the code =>

```
x = np.random.uniform(low=0, high=1, size=(num_train,))
```

the generating distribution of x is uniform.

- (2) Based upon the code =>

```
y = x - 2*x**2 + x**3 + np.random.normal(loc=0, scale=0.03, size=(num_train,))
```

the distribution of the additive noise (ϵ) is normal.

0.1.4 Fitting data to the model (5 points)

Here, we'll do linear regression to fit the parameters of a model $y = ax + b$.

```
[ ]: # xhat = (x, 1)
xhat = np.vstack((x, np.ones_like(x)))

# ===== #
# START YOUR CODE HERE #
# ===== #
# GOAL: create a variable theta; theta is a numpy array whose elements are [a,
↪ b]


theta = np.linalg.inv(xhat.dot(xhat.T)).dot(xhat.dot(y))

# ===== #
# END YOUR CODE HERE #
# ===== #
```

```
[ ]: # Plot the data and your model fit.
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')

# Plot the regression line
xs = np.linspace(min(x), max(x), 50)
xs = np.vstack((xs, np.ones_like(xs)))
plt.plot(xs[0,:], theta.dot(xs))
```

```
[ ]: [<matplotlib.lines.Line2D at 0x11ddd8250>]
```



0.1.5 QUESTIONS

- (1) Does the linear model under- or overfit the data?
- (2) How to change the model to improve the fitting?

0.1.6 ANSWERS

- (1) This linear model underfits the given data.
- (2) As we discussed in lecture, increasing the order of the polynomial used to characterize the model will lead to an improved fit of the data i.e. adding more polynomial terms e.g. cx^2

0.1.7 Fitting data to the model (5 points)

Here, we'll now do regression to polynomial models of orders 1 to 5. Note, the order 1 model is the linear model you prior fit.

```
[ ]: N = 5
      xhats = []
      thetas = []

      # ===== #
```

```

# START YOUR CODE HERE #
# ===== #

# GOAL: create a variable thetas.
# thetas is a list, where theta[i] are the model parameters for the polynomial
    ↪ fit of order i+1.
# i.e., thetas[0] is equivalent to theta above.
# i.e., thetas[1] should be a length 3 np.array with the coefficients of the
    ↪  $x^2$ ,  $x$ , and 1 respectively.
# ... etc.

for i in range(N):
    if i == 0:
        xhat = np.vstack((x, np.ones_like(x)))
        xhats.append(xhat)
        thetas.append(theta)
    else:
        xhat = np.vstack((x**(i + 1), xhat))
        xhats.append(xhat)
        thetas.append(np.linalg.inv(xhat.dot(xhat.T)).dot(xhat.dot(y)))

# ===== #
# END YOUR CODE HERE #
# ===== #

```

```

[ ]: # Plot the data
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')


# Plot the regression lines
plot_xs = []
for i in np.arange(N):
    if i == 0:
        plot_x = np.vstack((np.linspace(min(x), max(x), 50), np.ones(50)))
    else:
        plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))
    plot_xs.append(plot_x)

for i in np.arange(N):
    ax.plot(plot_xs[i][-2:], thetas[i].dot(plot_xs[i]))

labels = ['data']
[labels.append('n={}'.format(i+1)) for i in np.arange(N)]
bbox_to_anchor=(1.3, 1)

```

```
lgd = ax.legend(labels, bbox_to_anchor=bbox_to_anchor)
```



0.1.8 Calculating the training error (5 points)

Here, we'll now calculate the training error of polynomial models of orders 1 to 5.

```
[ ]: training_errors = []

# ===== #
# START YOUR CODE HERE #
# ===== #

# GOAL: create a variable training_errors, a list of 5 elements,
# where training_errors[i] are the training loss for the polynomial fit of
# order i+1.

# Calculate MSE

for i in range(N):
    theta = thetas[i]
    prediction = np.polyval(theta, x) #Find predicted value to compare against
    mse = np.mean((y - prediction)**2) * 100 #Calculate MSE
    training_errors.append(mse)

# ===== #
# END YOUR CODE HERE #
```

```
# ===== #  
  
print('Training errors are: \n', training_errors)
```

Training errors are:

```
[0.23799610883627006, 0.1092492220926853, 0.08169603801105368,  
0.08165353735296978, 0.08161479195525292]
```

0.1.9 QUESTIONS

- (1) What polynomial has the best training error?
- (2) Why is this expected?

0.1.10 ANSWERS


- (1) $n = 5$.
- (2) This is expected because it has higher degree polynomials to help better fit the data. As discussed in lecture, each polynomial degree increased should have AT LEAST just as good of a fit as the previous degree (because you could just put 0's for the higher degree coefficients).

0.1.11 Generating new samples and testing error (5 points)

Here, we'll now generate new samples and calculate testing error of polynomial models of orders 1 to 5.

```
[ ]: x = np.random.uniform(low=1, high=2, size=(num_train,))  
y = x - 2*x**2 + x**3 + np.random.normal(loc=0, scale=0.03, size=(num_train,))  
f = plt.figure()  
ax = f.gca()  
ax.plot(x, y, '.')  
ax.set_xlabel('$x$')  
ax.set_ylabel('$y$')
```

```
[ ]: Text(0, 0.5, '$y$')
```



```
[ ]: xhats = []
for i in np.arange(N):
    if i == 0:
        xhat = np.vstack((x, np.ones_like(x)))
        plot_x = np.vstack((np.linspace(min(x), max(x), 50), np.ones(50)))
    else:
        xhat = np.vstack((x**(i+1), xhat))
        plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))

    xhats.append(xhat)
```

```
[ ]: # Plot the data
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')

# Plot the regression lines
plot_xs = []
for i in np.arange(N):
```




```

if i == 0:
    plot_x = np.vstack((np.linspace(min(x), max(x), 50), np.ones(50)))
else:
    plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))
plot_xs.append(plot_x)

for i in np.arange(N):
    ax.plot(plot_xs[i][-2:], thetas[i].dot(plot_xs[i]))

labels = ['data']
[labels.append('n={}'.format(i+1)) for i in np.arange(N)]
bbox_to_anchor=(1.3, 1)
lgd = ax.legend(labels, bbox_to_anchor=bbox_to_anchor)

```



```

[ ]: testing_errors = []

# ===== #
# START YOUR CODE HERE #
# ===== #

# GOAL: create a variable testing_errors, a list of 5 elements,
# where testing_errors[i] are the testing loss for the polynomial fit of order
# i+1.

for i in range(N):

```

```

theta = thetas[i]
prediction = np.polyval(theta, x) #Find predicted value to compare against
mse = np.mean((y - prediction)**2) * 100 #Calculate MSE
testing_errors.append(mse)

# ===== #
# END YOUR CODE HERE #
# ===== #

print ('Testing errors are: \n', testing_errors)

```

Testing errors are:

```
[80.86165184550586, 213.1919244505791, 3.1256971084083736, 1.187076521149622,
214.91021747012803]
```

0.1.12 QUESTIONS

- (1) What polynomial has the best testing error?
- (2) Why polynomial models of orders 5 does not generalize well?

0.1.13 ANSWERS

- (1) $n = 4$.
- (2) This is due to overfitting as we discussed in lecture.

[]: