

```
import numpy as np
```

```
class Softmax(object):
```

```
    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)
```

```
    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001
```

```
    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
```

```
    # Initialize the loss to zero.
    loss = 0.0
```

```
    # ===== #
    # YOUR CODE HERE:
    #   Calculate the normalized softmax loss. Store it as the variable loss.
    #   (That is, calculate the sum of the losses of all the training
    #   set margins, and then normalize the loss by the number of
    #   training examples.)
    # ===== #
```

```
    shape = X.shape
    N = shape[0]
```

```
    for i in range(N):
        # Calculate score
        scores = X[i].dot(self.W.T)
        scores -= np.max(scores)

        # Softmax probabilities Computation
        softmax_probs = np.exp(scores) / np.sum(np.exp(scores))

        # Calculate the loss for the correct class
        correct_class = y[i]
        loss += -np.log(softmax_probs[correct_class])
```

```
    # Normalize the loss by the number of training examples
    loss /= N
```

```
    # ===== #
    # END YOUR CODE HERE
    # ===== #
```

```
    return loss
```

```

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
    the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and the gradient. Store the gradient
    # as the variable grad.
    # ===== #

    shape = X.shape
    N = shape[0]

    w_shape = self.W.shape
    K = w_shape[0]

    for i in range(N):
        # Calculate score
        scores = X[i].dot(self.W.T)
        scores -= np.max(scores)

        # Softmax probabilities Computation
        softmax_probs = np.exp(scores) / np.sum(np.exp(scores))

        # Calculate the loss for the correct class
        correct_class = y[i]
        loss += -np.log(softmax_probs[correct_class])

        # Gradient computation
        softmax_probs[correct_class] = (softmax_probs[correct_class] - 1)
        for j in range(K):
            grad[j, :] += softmax_probs[j] * X[i]

    # Normalize the loss by the number of training examples
    loss /= N
    grad /= N

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X, y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

```

```

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) +
abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic,
rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #

    shape = X.shape
    N = shape[0]

    # Calculate scores
    scores = X.dot(self.W.T)
    scores -= np.max(scores, axis=1, keepdims=True)

    # Calculate softmax prob.
    softmax_probs = np.exp(scores) / np.sum(np.exp(scores), axis=1, keepdims=True)

    # Loss calculation
    loss = (-np.sum(np.log(softmax_probs[np.arange(N), y]))) / N

    # Gradient computation
    softmax_probs[np.arange(N), y] = (softmax_probs[np.arange(N), y] - 1)
    grad = (softmax_probs.T.dot(X) / N)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
        batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

    # Run stochastic gradient descent to optimize W

```

```

loss_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    # Sample batch_size elements from the training data for use in
    # gradient descent. After sampling,
    # - X_batch should have shape: (batch_size, dim)
    # - y_batch should have shape: (batch_size,)
    # The indices should be randomly generated to reduce correlations
    # in the dataset. Use np.random.choice. It's okay to sample with
    # replacement.
    # ===== #

    shape = X.shape
    N = shape[0]

    index = np.random.choice(N, batch_size)
    X_batch = X[index]
    y_batch = y[index]

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # evaluate loss and gradient
    loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    # Update the parameters, self.W, with a gradient step
    # ===== #

    self.W = self.W - (grad * learning_rate)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    if verbose and it % 100 == 0:
        print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    # Predict the labels given the training data.
    # ===== #

    scores = X.dot(self.W.T)
    y_pred = np.argmax(scores, axis=1)

```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #  
  
return y_pred
```