

```

import numpy as np

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of  $D$ , a hidden dimension of  $H$ , and perform classification over  $C$  classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead, it
    will interact with a separate Solver object that is responsible for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
                 dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg

        # ===== #
        # YOUR CODE HERE:
        # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
        # self.params['W2'], self.params['b1'] and self.params['b2']. The
        # biases are initialized to zero and the weights are initialized
        # so that each parameter has mean 0 and standard deviation weight_scale.
        # The dimensions of W1 should be (input_dim, hidden_dim) and the
        # dimensions of W2 should be (hidden_dims, num_classes)
        # ===== #

        self.params['W1'] = np.random.normal(scale=weight_scale, size=(input_dim, hidden_dims))
        self.params['b1'] = np.zeros(hidden_dims)
        self.params['W2'] = np.random.normal(scale=weight_scale, size=(hidden_dims, num_classes))
        self.params['b2'] = np.zeros(num_classes)

        # ===== #
        # END YOUR CODE HERE
        # ===== #

    def loss(self, X, y=None):
        """
        Compute loss and gradient for a minibatch of data.

        Inputs:
        - X: Array of input data of shape (N, d_1, ..., d_k)
        - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

        Returns:

```

*If y is None, then run a test-time forward pass of the model and return:*

- *scores: Array of shape (N, C) giving classification scores, where scores[i, c] is the classification score for X[i] and class c.*

*If y is not None, then run a training-time forward and backward pass and return a tuple of:*

- *loss: Scalar value giving the loss*
- *grads: Dictionary with the same keys as self.params, mapping parameter names to gradients of the loss with respect to those parameters.*

```
"""
scores = None

# ===== #
# YOUR CODE HERE:
#   Implement the forward pass of the two-layer neural network. Store
#   the class scores as the variable 'scores'. Be sure to use the layers
#   you prior implemented.
# ===== #

scores, cache_hidden = affine_relu_forward(X.reshape(X.shape[0], -1), self.params['W1'],
self.params['b1'])
scores, cache_scores = affine_forward(scores, self.params['W2'], self.params['b2'])

# ===== #
# END YOUR CODE HERE
# ===== #

# If y is None then we are in test mode so just return scores
if y is None:
    return scores

loss, grads = 0, {}
# ===== #
# YOUR CODE HERE:
#   Implement the backward pass of the two-layer neural net. Store
#   the loss as the variable 'loss' and store the gradients in the
#   'grads' dictionary. For the grads dictionary, grads['W1'] holds
#   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
#   i.e., grads[k] holds the gradient for self.params[k].
#
#   Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
#   for each W. Be sure to include the 0.5 multiplying factor to
#   match our implementation.
#
#   And be sure to use the layers you prior implemented.
# ===== #

loss, dscores = softmax_loss(scores, y)
loss += 0.5 * self.reg * (np.sum(self.params['W1'] ** 2) + np.sum(self.params['W2'] ** 2))

# Compute gradients for the second layer
dH2, grads_W2, grads_b2 = affine_backward(dscores, cache_scores)
reg_W2 = self.reg * self.params['W2']
grads['W2'] = grads_W2 + reg_W2
grads['b2'] = grads_b2

# Compute gradients for the first layer
dX1, grads_W1, grads_b1 = affine_relu_backward(dH2, cache_hidden)
reg_W1 = self.reg * self.params['W1']
grads['W1'] = grads_W1 + reg_W1
grads['b1'] = grads_b1

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads
```

```
class FullyConnectedNet(object):
```

```
    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be
```

```
{affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
```

```
where batch normalization and dropout are optional, and the {...} block is
repeated L - 1 times.
```

```
Similar to the TwoLayerNet above, learnable parameters are stored in the
self.params dictionary and will be learned using the Solver class.
```

```
    """
```

```
def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=0, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):
```

```
    """
```

```
    Initialize a new FullyConnectedNet.
```

```
    Inputs:
```

- hidden\_dims: A list of integers giving the size of each hidden layer.
- input\_dim: An integer giving the size of the input.
- num\_classes: An integer giving the number of classes to classify.
- dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then the network should not use dropout at all.
- use\_batchnorm: Whether or not the network should use batch normalization.
- reg: Scalar giving L2 regularization strength.
- weight\_scale: Scalar giving the standard deviation for random initialization of the weights.
- dtype: A numpy datatype object; all computations will be performed using this datatype. float32 is faster but less accurate, so you should use float64 for numeric gradient checking.
- seed: If not None, then pass this random seed to the dropout layers. This will make the dropout layers deterministic so we can gradient check the model.

```
    """
```

```
self.use_batchnorm = use_batchnorm
self.use_dropout = dropout > 0
self.reg = reg
self.num_layers = 1 + len(hidden_dims)
self.dtype = dtype
self.params = {}
```

```
# ===== #
# YOUR CODE HERE:
# Initialize all parameters of the network in the self.params dictionary.
# The weights and biases of layer 1 are W1 and b1; and in general the
# weights and biases of layer i are Wi and bi. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation weight_scale.
# ===== #
```

```
layer_dims = [input_dim] + hidden_dims + [num_classes]
```

```
for i in range(self.num_layers):
    self.params["W{}".format(i + 1)] = weight_scale * np.random.randn(layer_dims[i],
layer_dims[i + 1])
    self.params["b{}".format(i + 1)] = np.zeros(layer_dims[i + 1])
```

```
# ===== #
# END YOUR CODE HERE
# ===== #
```

```

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
    if seed is not None:
        self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

```

```
def loss(self, X, y=None):
```

```

    """
    Compute loss and gradient for the fully-connected net.

```

```

    Input / output: Same as TwoLayerNet above.
    """

```

```

X = X.astype(self.dtype)
mode = 'test' if y is None else 'train'

```

```

# Set train/test mode for batchnorm params and dropout param since they
# behave differently during training and testing.

```

```

if self.dropout_param is not None:
    self.dropout_param['mode'] = mode

```

```

if self.use_batchnorm:
    for bn_param in self.bn_params:
        bn_param[mode] = mode

```

```
scores = None
```

```

# ===== #
# YOUR CODE HERE:
#   Implement the forward pass of the FC net and store the output
#   scores as the variable "scores".
# ===== #

```

```

caches = {}

```

```

for i in range(1, self.num_layers + 1):
    W = self.params[f"W{i}"]
    b = self.params[f"b{i}"]

```

```

# Forward pass through layers with ReLU activation

```

```

X, cache = affine_relu_forward(X, W, b) if i < self.num_layers else affine_forward(X, W,

```

```

    caches[i] = cache

```

```
scores = X
```

```

# ===== #
# END YOUR CODE HERE
# ===== #

```

b)

```

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
#   Implement the backwards pass of the FC net and store the gradients
#   in the grads dict, so that grads[k] is the gradient of self.params[k]
#   Be sure your L2 regularization includes a 0.5 factor.
# ===== #

# Compute the softmax loss and its gradient
loss, dout = softmax_loss(scores, y)

# Add regularization to the loss more efficiently
weights_squared_sum = sum(np.sum(self.params['W' + str(i + 1)] ** 2) for i in
range(self.num_layers))
loss += 0.5 * self.reg * weights_squared_sum

# Backward pass for the last layer
dout, grads['W' + str(self.num_layers)], grads['b' + str(self.num_layers)] =
affine_backward(dout, caches[self.num_layers])
grads['W' + str(self.num_layers)] += self.reg * self.params['W' + str(self.num_layers)]

# Iterate over layers in reverse for backprop, skipping the last since it's already handled
for i in range(self.num_layers - 1, 0, -1):
    dout, dW, db = affine_relu_backward(dout, caches[i])
    grads['W' + str(i)] = dW + self.reg * self.params['W' + str(i)]
    grads['b' + str(i)] = db

# ===== #
# END YOUR CODE HERE
# ===== #
return loss, grads

```